Reproducing PIM-zd-tree on a Simulated UPMEM System with Skew-Aware Point Lookups

Chetan Srirama
Roll no. 240101031

Samrudhhi Sonawane
Roll no. 240101083

Venkata Sai Shankar Karthik Macharla
Roll no. 240101111

***Abstract*—PIM-zd-tree (PPoPP'26) is a spatial index for processing-in-memory (PIM) systems that splits a zd-tree into layers with different placement and caching rules, and reports up to $4.25\times$ higher throughput than shared-memory indexes on a 2560-core UPMEM server. We reproduce its core claims at reduced scale: the authors' public code runs on the UPMEM SDK functional simulator with 8–64 simulated PIM cores (DPUs), and performance comes from DPU cycle counts, counted transfer bytes, transfer bandwidths measured by PrIM, and measured host time. Before the code ran correctly at this scale we had to fix memory-layout overflows and a host-side bug that lost points under skewed data. We then compare PIM-zd-tree against a range-partitioned PIM design and the Pkd-tree and zd-tree on the CPU, for uniform and skewed data. The paper's trends reproduce: throughput grows with the number of DPUs (lookups $5\times$ from 8 to 64 DPUs, overtaking the 20-thread Pkd-tree), placement by data size keeps the slowest DPU within 1.06–$1.32\times$ of the mean where range partitioning reaches $25\times$, and inserts are host-bound. The paper's insert speedup over CPU baselines does not appear at 64 DPUs. As our addition we target a limitation the authors acknowledge: the throughput-optimized configuration degrades under skewed queries. Host-side lookup deduplication, a small hot-key cache and a per-batch adaptive switch make point lookups 4.3–$58\times$ faster under Zipf-skewed keys on 64 DPUs, and leave uniform workloads unchanged.**

***Index Terms*—processing-in-memory, UPMEM, spatial index, zd-tree, load balance, skew**

I. Introduction

Space-partitioning indexes such as kd-trees and quad/octrees answer point, range and nearest-neighbour queries over multi-dimensional data, and in-memory versions are usually limited by memory bandwidth. Processing-in-memory (PIM) attaches a simple core to each memory bank, so work moves to the data instead of data moving to the CPU. The UPMEM system is the first commercial PIM product: thousands of in-memory cores (DPUs), each with its own 64 MB bank (MRAM), connected only through the host CPU.

PIM-zd-tree [1] adapts the zd-tree [2] to this setting. The tree is divided by subtree size into three layers: a globally shared top layer (L0) kept on the CPU, partially shared L1 nodes placed on random PIM modules with cached ancestors and descendants, and exclusive L2 nodes. Two thresholds, θ_{L0} and θ_{L1}, trade skew tolerance against communication. The paper evaluates a *throughput-optimized* configuration ($\theta_{L0} = n/P$, $\theta_{L1} = 1$) on a 2560-DPU server, and a *skew-resistant* one in its skew study (Fig. 9).

The assignment asks us to (i) reproduce the core idea at reduced scale with the authors' public code on a simulator, running batched inserts and point lookups on 8–64 simulated PIM cores for uniform and skewed data, comparing throughput trends and per-core load balance against a shared-memory CPU index; and (ii) add something of our own and evaluate it against the baseline. Our contributions:

- A working, documented port of the public code to the UPMEM functional simulator, with per-DPU cycle and byte counters, a reproducible data generator calibrated to the paper's two real data sets, and fixes for three overflows and one correctness bug that only show up at small DPU counts or under skew (Sec. IV).
- A reduced-scale evaluation against a range-partitioned PIM design and the Pkd-tree and zd-tree on the CPU, compared trend by trend with the paper (Sec. V).
- Skew-aware point lookups for the throughput-optimized configuration: deduplication, a hot-key cache and adaptive switching, evaluated across query skew, cache size and DPU count (Sec. VI).

The code is at https://github.com/⟨your-group⟩/PIM-zd-tree/tree/simulator.

II. Background

PIM model and UPMEM. In the PIM model [3] a host CPU with a cache of M words is connected to P PIM modules, each holding $\Theta(N/P)$ words. Modules cannot talk to each other; programs run in bulk-synchronous rounds, and a round lasts as long as its slowest module, so load balance across modules decides performance. On UPMEM, each DPU runs up to 24 hardware threads (tasklets) at 350 MHz on its private 64 MB MRAM. Host–DPU transfers are fastest when every DPU receives a buffer of the same size: a parallel transfer pads every buffer to the largest one [4].

zd-tree. A zd-tree [2] is a compressed binary trie over the Morton (z-order) keys of the points: the root splits on the first key bit, single-child chains are compressed, and leaves hold at most 16 points. It is deterministic (history-independent) and needs no rebalancing, which suits PIM.

PIM-zd-tree. Nodes with subtree size $T \geq \theta_{L0}$ form L0 and are replicated on the CPU. Nodes with $\theta_{L1} \leq T < \theta_{L0}$ form L1: each sits on a random module, together with cached copies of its ancestors and descendants. Smaller nodes form L2 and are not replicated. Lazy counters keep approximate subtree sizes; push-pull search moves heavily accessed meta-nodes to the CPU. In the throughput-optimized configuration every subtree below L0 has about n/P points and lives on

TABLE I
SCOPE OF THE REPRODUCTION

Aspect	Paper	This work
PIM hardware	UPMEM, 2560 DPUs	UPMEM SDK 2025.1.0 functional simulator, 8–64 DPUs (one rank)
PIM time	measured	simulator cycle counters / 350 MHz
Transfers	measured, direct interface	bytes counted exactly, time from PrIM bandwidths [4]
Host CPU	2×Xeon 4216	i5-13500, 20 threads (measured)
CPU baselines	Pkd-tree, zd-tree, 40-thread Xeon	same trees (SpaceTreeLib [5]), 1/4/20 threads
Data size	300M points	0.8M–1M points (≈119k per DPU at 8 DPUs)
Data sets	uniform, COSMOS, OSM	uniform; Zipf clusters with Gini 0.28 and 0.98
Operations	insert, box count/fetch, kNN	insert, point lookup
Variants	throughput-opt., skew-resistant	throughput-opt. (public code)
Memory traffic	PAPI counters	CPU–PIM bytes only

one module, so a search costs one round trip. The public code implements this configuration only: a fixed key-range split on the CPU and one local subtree per DPU, with no push-pull and no L1 caching.

III. Methodology

A. What we kept, left out and replaced

All real-hardware experiments are replaced by simulation and a cost model. Table I lists the differences from the paper.

B. Simulator and cost model

We use the UPMEM SDK 2025.1.0 functional simulator, whose API headers match the SDK fork the authors build against. uPIMulator [6] is cycle-level but cannot run the authors' C++ host program, so we did not use it. The simulator's wall-clock time says nothing about real hardware. We instead record, per batch, each DPU's cycle counter for every launch (reset by tasklet 0 at launch start, read after all tasklets finish), the bytes moved to and from every DPU, and the host's own wall time with simulator time subtracted. The modelled batch time is

$$T = T_{\text{host}} + \sum_{\text{launches}} \frac{\max_i c_i}{350\,\text{MHz}} + \frac{B_{\rightarrow}}{BW_{\rightarrow}(n)} + \frac{B_{\leftarrow}}{BW_{\leftarrow}(n)}, \quad (1)$$

where c_i are DPU i's cycles, B are the padded bytes actually transferred, and $BW(n)$ is a power-law fit through PrIM's measured one-rank bandwidths [4]: CPU→DPU 0.33 GB/s (1 DPU) to 6.68 GB/s (64 DPUs), DPU→CPU 0.12 to 4.74 GB/s. We assume no launch overhead. Load imbalance is the slowest DPU's cycles divided by the mean. Fig. 2 shows the setup.

C. Data, workloads and baselines

Points are 3-D with 31-bit coordinates. The *skewed* data sets place 1024 cluster centres uniformly, choose a cluster with Zipf probability $\propto 1/r^s$ and scatter points around its centre. Following the paper we measure skew as the Gini coefficient of point counts over 2048 equal key ranges, and calibrate a *moderate* set (Gini 0.28; paper's COSMOS 0.287) and a *strong* set (Gini 0.98; OSM 0.967). Every point is a pure function of (seed, batch, index), so runs are reproducible.

Each run inserts 800k warm-up points in 10 batches, then measures three batches of 50k inserts and 50k point lookups. We compare four designs: *PIM-zd-tree* with each DPU's key range holding an equal share of the warm-up points (the subtree-size split, $\theta_{L0} = n/P$); a *range-partitioned PIM* design with equal-width key ranges (the public code's default, which the paper's Sec. 3 argues against); and the *Pkd-tree* [7] and *zd-tree* [2] from SpaceTreeLib on the CPU. A verification mode looks up every inserted point after every batch, and every reported run returned the correct answer for every lookup.

IV. Running the public code at small scale

The public code targets 2560 DPUs, where each DPU receives a small slice of every batch. At 8–64 DPUs a DPU gets 40–300× more, and several fixed-size limits were exceeded silently, because the code's assertions are compiled out in normal builds.

Simulator setup. We allocate exactly `NR_DPUS` ("all" gives one 64-DPU rank on the simulator), use the standard UPMEM transfer interface (the direct interface needs real ranks and AVX-512), and build one binary set per DPU count.

Overflows. (i) An insert task carrying more than 3072 points overran its tasklet's 128 KB MRAM scratch area; on an empty tree each DPU's whole first batch lands in one or two tasks. (ii) More than 12,800 insert tasks per DPU overran a host offset array. (iii) The DPU reply area was only 0.93 MB, about 119k replies. We cap task length and tasks per DPU and defer the rest to an extra round, rebalance MRAM so replies get 1.9 MB, and add always-on capacity checks that report the cause instead of faulting.

Lost points under skew. With clustered data, lookups failed for points inserted in *earlier* batches. The host groups points into one task per target by scanning key order. New points on both sides of a dense subtree C target the same compressed edge above C, but C's own descendants sit between them in key order, so the host sent two tasks for one edge. The second edge split built a node inside the first one's prefix and disconnected C (27 of 80k points lost in one test). Grouping by target before building tasks fixes it, matching the paper's statement that the CPU deduplicates conflicting inserts (Alg. 2).

V. Baseline results

Scaling with PIM cores (Fig. 1). From 8 to 64 DPUs, PIM-zd-tree's throughput rises 4.0–4.5× for inserts and 5.0–5.5× for lookups (ideal: 8×). Lookups overtake the 20-thread Pkd-tree at 32–64 DPUs (22.3 vs. 15.0 Mop/s, uniform); inserts reach 7.5 Mop/s against the Pkd-tree's 18.3. The paper reports PIM-zd-tree 1.82× faster than the Pkd-tree on inserts at 2560 DPUs. We reproduce the direction but not the crossover: 64

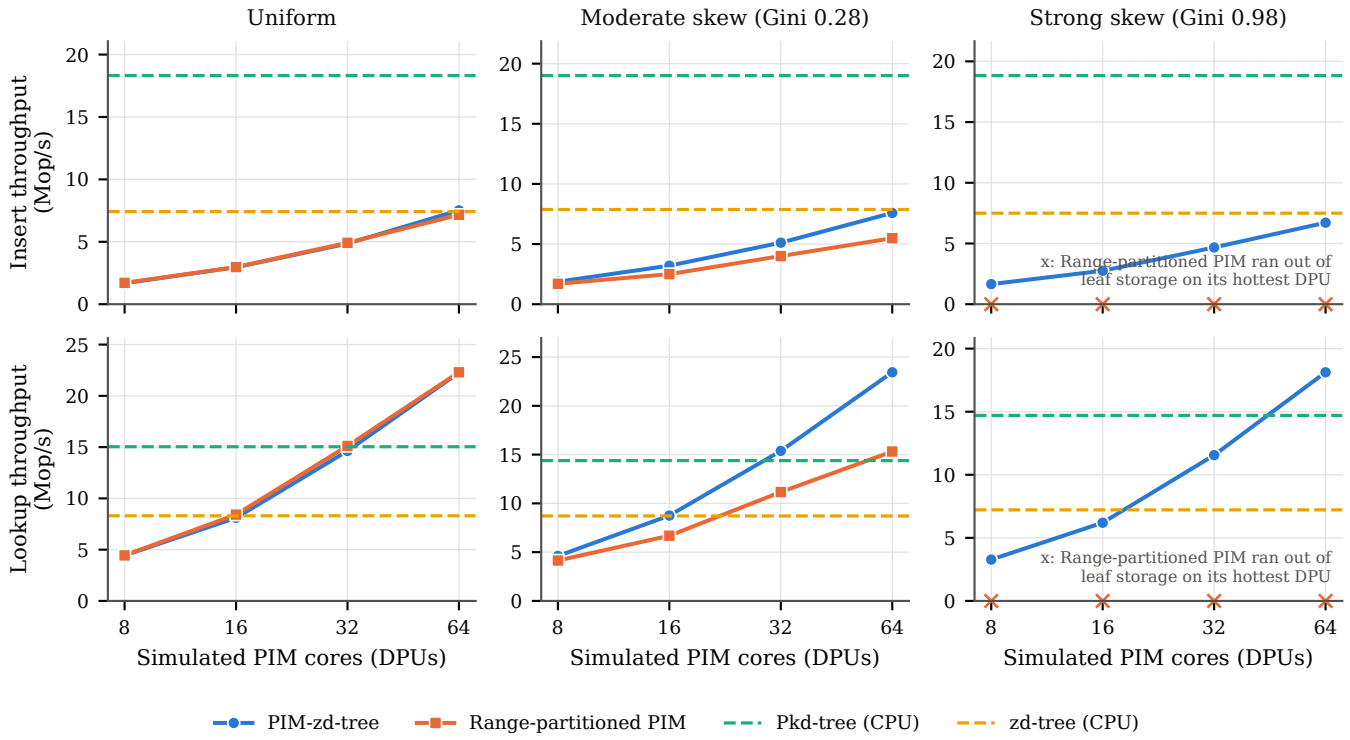


Fig. 1. Insert (top) and lookup (bottom) throughput vs. number of simulated DPUs, 800k warm-up points. Range-partitioned PIM cannot hold the strongly skewed data set: its hottest DPU runs out of leaf storage at every size.

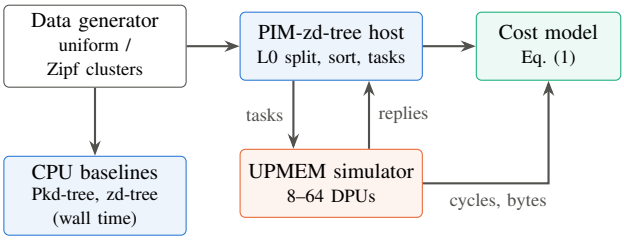


Fig. 2. Experimental setup. The cost model combines the host's measured time with the simulator's DPU cycles and transferred bytes. The PIM host and the CPU baselines get identical points and queries from one shared generator.

DPUs are 2.5% of the paper's PIM cores and bandwidth, and our CPU is a recent desktop part.

Load balance under skew (Fig. 3). This is the paper's central argument for its layout, and it reproduces clearly. Range partitioning puts the dense clusters on a few DPUs: with moderate skew its lookup imbalance grows to 2.1× at 64 DPUs (15.3 vs. 23.4 Mop/s for PIM-zd-tree). With strong skew the 800k-point set overflows its hottest DPU at every size. At 200k points it fits, but throughput stays flat at 1.0 Mop/s (lookup) and 0.6 Mop/s (insert) from 8 to 64 DPUs, because imbalance grows from 3.4× to 24.7×: extra DPUs do not help when one DPU does the work. Splitting by data size keeps PIM-zd-tree within 1.06–1.32×, and strong skew costs it only 19% of its uniform lookup throughput at 64 DPUs.

Where the time goes (Fig. 6). As in the paper's Fig. 6, inserts are the most host-heavy operation (sorting, grouping and task packing: 42–45% at 64 DPUs), while lookups are PIM-dominated (61–66%). Transfers take 6–8%. Host time does not shrink as DPUs are added, so the host's share of insert time grows from about 11% at 8 DPUs to 44% at 64. That host work is what limits insert scaling at our size.

Batch size (Fig. 4). As in the paper's Fig. 7, larger batches are faster: from 5k to 200k operations, inserts rise from 4.3 to 10.1 Mop/s and lookups from 12.7 to 29.9 Mop/s, as fixed per-round costs are spread out and per-DPU work evens out. The paper's drop for very large batches comes from L3 misses on the host, which our batch sizes do not reach.

Communication. PIM-zd-tree moves 58 B per insert and 16 B per lookup that are actually needed. Padding to the largest DPU's buffer adds 2–9%, growing with the DPU count.

VI. Our addition: skew-aware point lookups

A. Problem

The throughput-optimized configuration tolerates only mild skew. Every lookup goes to the module owning its key's subtree, and the paper's Fig. 9 shows it slowing by up to 10.66× when 2% of queries come from a skewed distribution. The cure the paper offers is the skew-resistant configuration, which is slower without skew and is not in the public code. Our setup shows the same problem: with Zipf-distributed lookup keys over 800k uniform points on 64 DPUs, the busiest

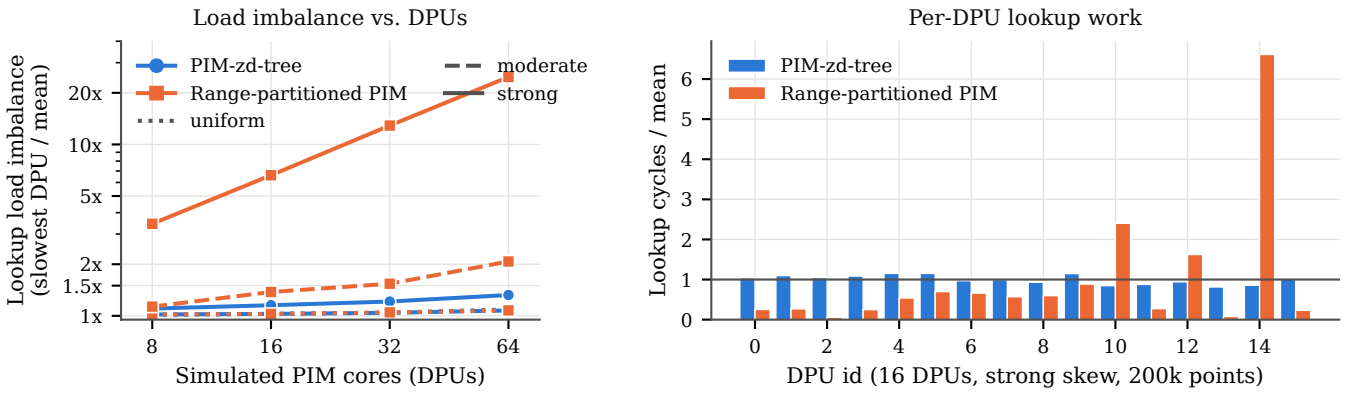


Fig. 3. Left: lookup load imbalance (slowest DPU / mean) vs. DPUs (strong skew from a 200k-point run that both designs can hold). Right: per-DPU lookup work at 16 DPUs under strong skew.

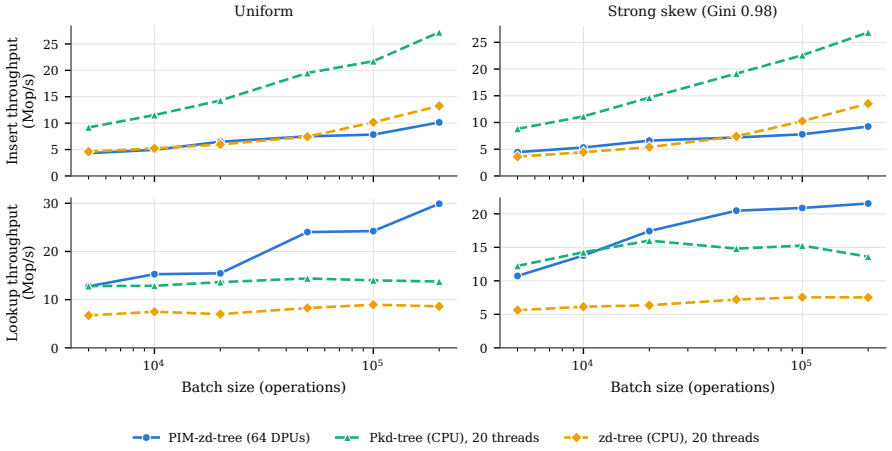


Fig. 4. Throughput vs. batch size at 64 DPUs; CPU baselines with 20 threads (median of 3 runs).

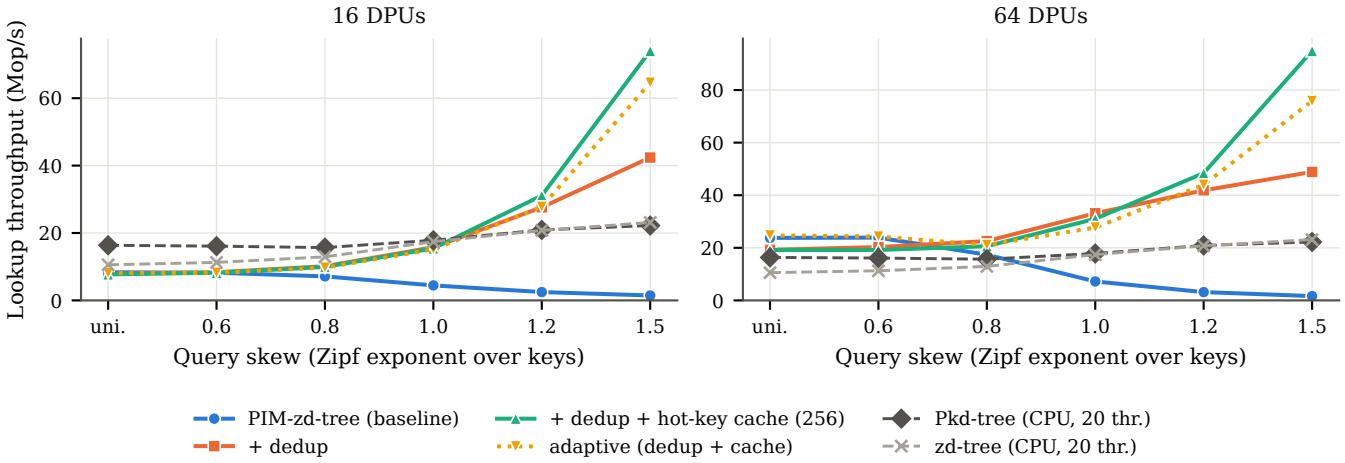


Fig. 5. Lookup throughput vs. query skew (Zipf exponent over keys), 10 batches of 50k lookups after 800k uniform points (median batch). CPU baselines see the same query stream.

DPU does 4.4× (s=1.0) to 21× (s=1.5) the mean work, and lookups fall from 23.8 to 1.6 Mop/s, below both CPU indexes.

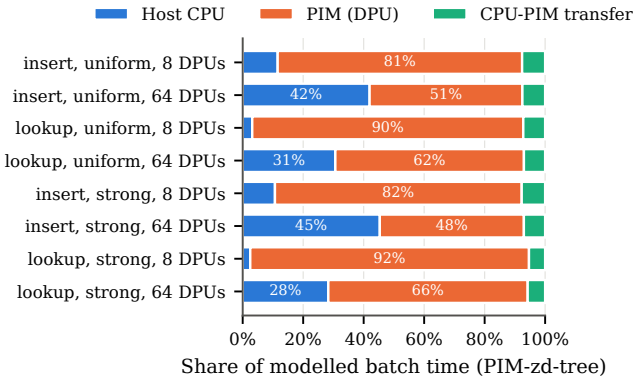


Fig. 6. Share of modelled batch time for PIM-zd-tree.

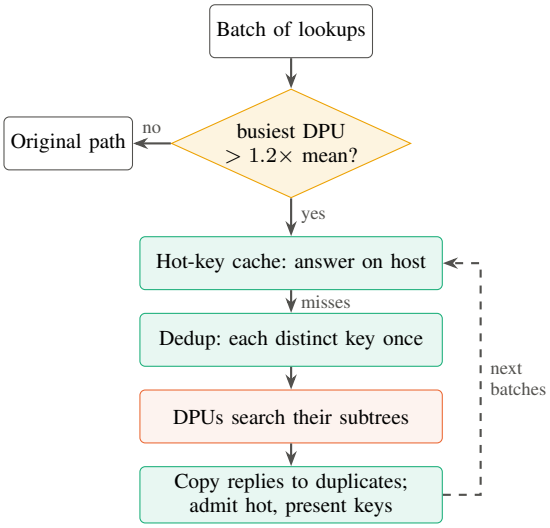


Fig. 7. Skew-aware lookup path (new parts in green).

B. Design

Fig. 7 shows the new lookup path. It runs on the host only; the DPU program is unchanged.

Deduplication. Each distinct key in a batch is sent once, and its reply is copied to every duplicate. This removes the BSP straggler that a hot key creates. Distinct keys are found with a lock-free open-addressing hash table kept between batches; for 50k keys it takes 0.2 ms, about 10× faster than sorting.

Hot-key cache. A small host-side set (256 keys) of keys that recent batches showed to be repeated and present. Their lookups are answered on the CPU. Entries cannot go stale: a lookup returns the key itself exactly when the key is stored, and the workload has no deletes (a delete would also remove the key from the cache). Admission uses exponentially decayed in-batch counts, restricted to the batch's most repeated keys so maintenance stays bounded.

Adaptive switching. Both mechanisms cost host time and pay off only when a round would wait on a hot DPU. The host knows every lookup's target DPU before sending, so it builds a per-DPU histogram and takes the new path only if the busiest DPU would get more than 1.2× the mean. The original path needs those targets anyway. This is the paper's push-pull idea (pull hot work to the CPU) applied per key and across batches, on the unchanged throughput-optimized layout.

TABLE II
LOOKUP THROUGHPUT AT 64 DPUS (MOP/S)

Skew	Baseline	+dedup	+dedup+cache	adaptive	Pkd (CPU)
uniform	23.8	19.3	19.2	**24.7**	16.4
s=0.6	23.9	20.3	19.2	**24.3**	16.1
s=0.8	17.4	**22.6**	20.6	21.1	15.7
s=1.0	7.2	**33.1**	31.0	27.7	17.8
s=1.2	3.2	41.9	**48.4**	44.0	20.9
s=1.5	1.6	48.9	**94.9**	75.8	22.2

C. Evaluation

Queries are stored points whose popularity rank follows Zipf(s); ranks are scattered over the data by a fixed permutation. Throughput is batch size divided by the median modelled batch time over 10 batches, because host time is measured in milliseconds and the median keeps one-time start-up costs out. All 56 PIM runs returned the correct answer for every lookup.

Speedup (Fig. 5, Table II). Under skew, the new path is 1.2–1.3× faster at s=0.8, 4.3–4.6× at s=1.0, 13–15× at s=1.2 and 30–58× at s=1.5 on 64 DPUs; 16 DPUs show the same pattern (1.4–51×). With the new path, PIM-zd-tree beats both CPU indexes from s=0.8, where the baseline falls below them. With strongly skewed *data* and hot keys together (s=1.0), dedup raises lookups from 5.7 to 29.6 Mop/s.

Why it works (Fig. 8). Imbalance stays at 1.10–1.23× at every skew, against 21× for the baseline. At s=1.5 only 4% of lookups still reach the DPUs.

Cache size and cost (Fig. 9). Dedup alone handles moderate skew; the cache pays off at heavy skew, nearly doubling throughput at s=1.5 (94.9 vs. 48.9 Mop/s) by removing the remaining hot keys from every batch. At s=1.2, 64–256 keys are best (53 and 48 Mop/s). Larger caches add host maintenance without catching more hot queries. Neither mechanism is free: at 64 DPUs a uniform batch takes about 2 ms, and dedup plus cache adds about 0.5 ms of host work, a 19% loss without skew. The adaptive switch removes that loss: it took the original path in all 10 batches for uniform and s=0.6 queries (24.7 and 24.3 vs. 23.8 and 23.9 Mop/s) and the new path in every batch from s=0.8. It trails the best fixed variant at extreme skew by the histogram's cost. With the new path the batch time at s=1.2 falls from 15.8 ms (89% PIM) to 1.0 ms, mostly host time: the DPUs are no longer the bottleneck.

VII. LIMITATIONS AND THREATS TO VALIDITY

Model, not hardware. PIM time comes from the functional simulator's cycle counts, and transfer time from published bandwidths with no launch overhead, so small batches may look better than on real hardware. Comparisons between designs on the same simulator are more reliable than absolute numbers. **Scale.** 64 DPUs are 2.5% of the paper's machine. The insert crossover with CPU indexes is not reached, and

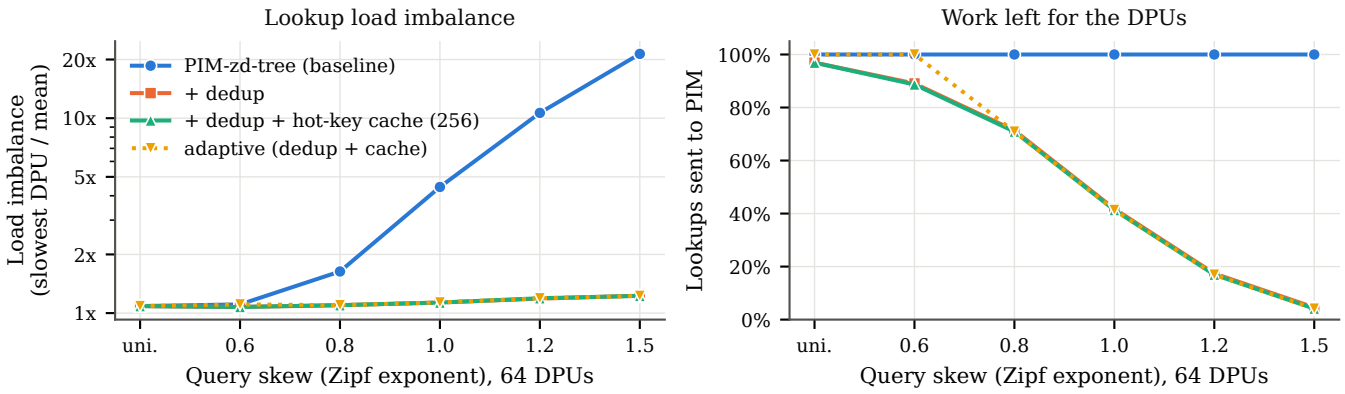


Fig. 8. 64 DPUs. Left: lookup load imbalance. Right: fraction of lookups still sent to the DPUs.

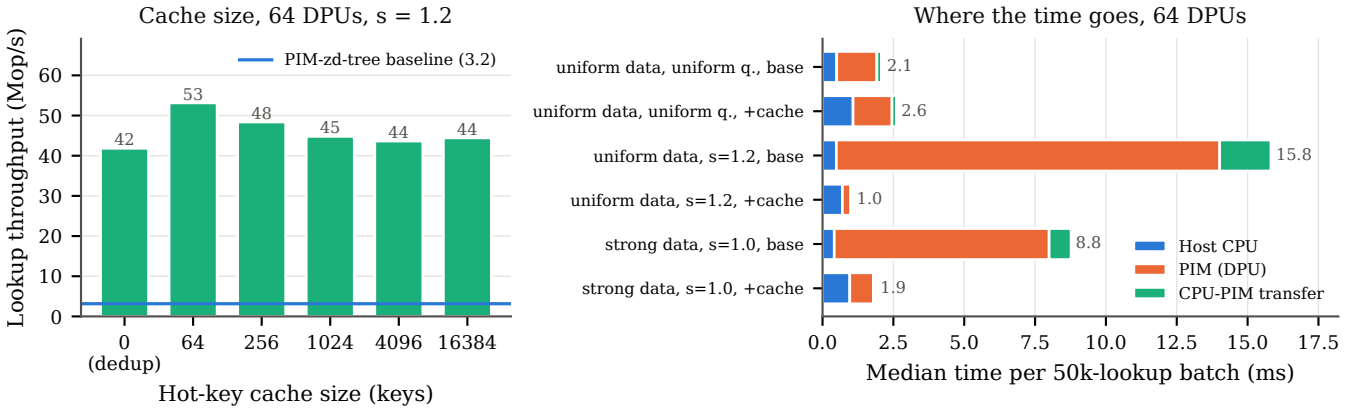


Fig. 9. Left: cache size at s=1.2. Right: median modelled time per batch, split by component.

the host share of time is higher than on the paper's server. **Public code.** It lacks push-pull search, L1 caching and the skew-resistant configuration, and our L0 borders are fixed from the warm-up data rather than adapted as data arrives. **Scope of the addition.** It handles hot *keys*. Many distinct keys in one small region still overload one DPU and need re-partitioning. It covers point lookups only, and the cache assumes no deletes.

VIII. Conclusion

At 8–64 simulated DPUs the authors' code reproduces PIM-zd-tree's main qualitative results: throughput grows with PIM cores, placement by subtree size keeps DPUs balanced where range partitioning collapses under skew, inserts are host-bound and queries PIM-bound, and larger batches help. The insert advantage over CPU indexes needs far more DPUs than we could simulate. Getting there took fixes for buffer overflows and a host bug that lost points under skewed data. Our addition, deduplicated and cached point lookups with a per-batch adaptive switch, removes the throughput-optimized configuration's weakness under skewed queries (4.3–58× at 64 DPUs) without the skew-resistant layout and without slowing uniform workloads.

References

[1] Y. Zhao, H. Kang, Z. Men, Y. Gu, G. E. Blelloch, L. Dhulipala, C. McGuffey, and P. B. Gibbons, "PIM-zd-tree: A fast space-partitioning index leveraging processing-in-memory," in *Proceedings of the 31st ACM SIGPLAN Annual Symposium on Principles and Practice of Parallel Programming (PPoPP '26)*, pp. 480–495, 2026.

[2] G. E. Blelloch and M. Dobson, "Parallel nearest neighbors in low dimensions with batch updates," in *Proceedings of the Symposium on Algorithm Engineering and Experiments (ALENEX)*, 2022.

[3] H. Kang, P. B. Gibbons, G. E. Blelloch, L. Dhulipala, Y. Gu, and C. McGuffey, "The processing-in-memory model," in *Proceedings of the 33rd ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, 2021.

[4] J. Gómez-Luna, I. El Hajj, I. Fernandez, C. Giannoula, G. F. Oliveira, and O. Mutlu, "Benchmarking a new paradigm: Experimental analysis and characterization of a real processing-in-memory system," *IEEE Access*, vol. 10, pp. 52565–52608, 2022.

[5] Z. Men, B. Huang, Y. Gu, and Y. Sun, "Parallel dynamic spatial indexes," in *Proceedings of the 31st ACM SIGPLAN Annual Symposium on Principles and Practice of Parallel Programming (PPoPP '26)*, 2026.

[6] B. Hyun, T. Kim, D. Lee, and M. Rhu, "Pathfinding future PIM architectures by demystifying a commercial PIM technology," in *Proceedings of the IEEE International Symposium on High-Performance Computer Architecture (HPCA)*, 2024.

[7] Z. Men, Z. Shen, Y. Gu, and Y. Sun, "Parallel kd-tree with batch updates," *Proceedings of the ACM on Management of Data*, vol. 3, no. 1, pp. 1–26, 2025.