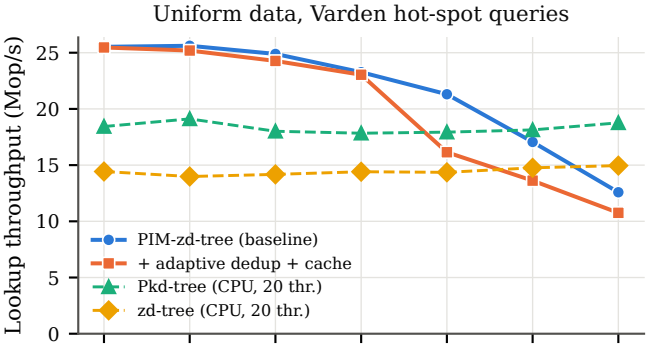
Uniform data, Varden hot-spot queries
Lookup throughput (Mop/s)
25
20
15
10
5
0
PIM-zd-tree (baseline)
+ adaptive dedup + cache
Pkd-tree (CPU, 20 thr.)
zd-tree (CPU, 20 thr.)
0
0.1
1
10
20
50
100
Share of lookups from Varden points (‰), 64 DPUs

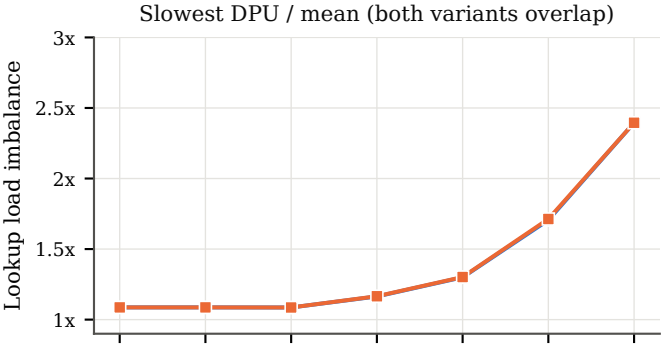
Slowest DPU / mean (both variants overlap)
Lookup load imbalance
3x
2.5x
2x
1.5x
1x
0
0.1
1
10
20
50
100
Share of lookups from Varden points (‰), 64 DPUs